

The Sleeping Teaching Assistant

Statement of the Problem and the Solution

(1)Overview:

In this tutorial, we'll explain the Sleeping Teaching Assistant.

It's a famous inter-process communication and synchronization problem that takes place in a TA 's office.

(2) Problem Definition

the problem takes place in a TA 's office. A university computer science department has a teaching assistant (TA) who helps undergraduate students with their programming assignments during regular office hours. The TA 's office is rather small and has room for only one desk with a chair and computer. There are three chairs in the hallway outside the office where students can sit and wait if the TA is currently helping another student. When there are no students who need help during office hours, the TA sits at the desk and takes a nap. If a student arrives during office hours and finds the TA sleeping, the student must awaken the TA to ask for help. If a student arrives and finds the TA currently helping another student, the student sits on one of the chairs in the hallway and waits. If no chairs are available, the student will come back at a later time.

- Let's look at the characteristics of the problem:

- We need to synchronize the Teaching Assistant and the students so there wouldn't be any race conditions. Since we have limited resources and waiting tasks, this problem has so many similarities with various queueing situations.
- students can enter the waiting hallway and should wait for whether the teacher is available or not
- If other students keep coming while the teacher is helping a student, they sit down in the waiting room
- students leave if there is no empty chair in the hallway

(3) Solution:

Our solution uses three semaphores, students, which counts waiting students(excluding the students in the teacher desks, who is not waiting), teacher, the number of teachers (0 or 1) who are sleeping, waiting for students, and mutex, which is used for mutual exclusion. We also need a variable, waiting, which also counts the waiting students. The reason for having waiting is that there is no way to read the current value of a semaphore. In this solution, a student entering the hallway has to count the number of waiting students. If it is less than the number of chairs, he stays; otherwise, he leaves.

- We can use the following variables for synchronisation and mutual exclusion and 2 threads .

```
#define CHAIRS 5
typedef int semaphore;
semaphore customers = 0;
semaphore teacher= 0;
semaphore mutex = 1;
int waiting = 0;
void teachers(void) {
    while (TRUE) {
        down(&students);
        down(&mutex);
        waiting = waiting - 1;
        up(&teachers);
        up(&mutex);
        Helps();
    }
}
void customer(void) {
    down(&mutex);
    if (waiting < CHAIRS) {
        waiting = waiting + 1;
        up(&students);
        down(&teacher);
        get_Help();
    }
    else {
```

```
        up(&mutex);  
    } }
```

```
void students(void)  
{ down(&mutex);  
  if (waiting < CHAIRS) {  
      waiting = waiting + 1;  
      up(&students);  
      up(&mutex);  
      down(&teachers);  
      get_Help();  
  } else { up(&mutex); } }
```

4. Conclusion:

In this tutorial, we've given a brief definition of the problem and then shared the solution. It's another important problem in concurrent programming because it applies to other types of queuing problems in computing, networks, industrial engineering, telecommunication, and traffic engineering.